



[View my portfolio](#)

CONTACT

- New York, USA
- rajani.maski@gmail.com
- (516) 309-0398
- [linkedin.com/in/rajanimaski](https://www.linkedin.com/in/rajanimaski)

SPEAKING

Optimized AI Conference 2026

Agentic Search Patterns **SPEAKER**

OpenSearchCon India 2026

Generative Discovery: Intent, Context and Cognition **SPEAKER**

CERTIFICATIONS

AWS Certified GenAI Developer Professional

DeepLearning.AI

40+ courses — RAG, Agents, LLMOps, Transformers

EDUCATION

M.S. Computer Science

San Francisco State University

B.E. Engineering

Visvesvaraya Technological University

CORE AI SKILLS

Agentic Systems

Multimodal AI

Prompt Engineering

RAG Pipelines

Vector Search

Embeddings

LLMOps

OpenSearch

AWS Bedrock

NVIDIA NIM

NVIDIA NeMo

NeMo Retriever

Triton Inference

cuVS / RAPIDS

TensorRT-LLM

LANGUAGES & INFRA

Python

Java

Scala

Bash

AWS

Kubernetes

Terraform

Rajani Maski

Staff Software Engineer, AI

PROFILE

Staff AI Engineer with 17+ years of experience from classical IR to modern AI, including model training, semantic search, hybrid retrieval, relevance ranking, learning-to-rank (LTR), and production retrieval systems serving 2000+ RPS. Brings hands-on experience working directly on customer projects as a Senior Consulting Engineer, and now focuses on Generative Discovery, building agentic and multimodal systems across retrieval, ranking, and image generation. Operating at the intersection of retrieval, intelligence, and real-world customer impact.

EXPERIENCE

Staff Software Engineer, AI

[Shutterstock](#) · New York · May 2022 – Present · 4 yrs 1 mo

- Architecting **Generative Discovery infrastructure** powering RAG-style GenAI applications, combining dense vector retrieval, hybrid search, and metadata grounding at 2000+ RPS across hundreds of nodes.
- Drove **25% latency and cost reduction** through retrieval architecture optimization while simultaneously improving relevance metrics, decoupling scale from quality tradeoffs.
- Building **multimodal embedding pipelines** supporting image + text semantic search, with MLflow and SageMaker for feature generation and model lifecycle management.
- Continuously improving the **ranking pipeline** to retrieve and rerank results for the best candidate match, leveraging Learning-to-Rank (LTR), bi-encoders, cross-encoders, and related techniques.
- Working with the team to automate **asset metadata tag generation** pipelines for images and videos, improving search accuracy and retrieval coverage.
- Designed **offline and online evaluation frameworks** for retrieval quality using precision/recall, relevance judgments, and human-in-the-loop feedback loops.
- Led **Solr to OpenSearch migration**, redesigning the retrieval layer for cloud-native scale and enabling ML-powered ranking pipelines including Learning-to-Rank.

Professional Services Consulting Engineer

[Redis](#) · New York · Dec 2020 – May 2022 · 1 yr 5 mos

- Designed and deployed **vector similarity search systems** using Redis as a senior consulting engineer, advising enterprise customers and translating business discovery needs into production retrieval architectures and reusable reference designs.
- Delivered projects using **RediSearch and RedisGraph** for financial banking customers, building knowledge graphs for CDC of banking transactions.

Senior Search Consultant

[Lucidworks](#) · New York · Jun 2018 – Dec 2020 · 2 yrs 6 mos

- Architected enterprise search platforms using **Lucidworks Fusion, Solr and Spark**; implemented collaborative recommendations, personalized search, and signal-based boosting with ML-based relevance pipelines.
- Led customer-facing relevance engagements end-to-end — from discovery and signal modeling to LTR deployment — measurably improving search quality for enterprise clients.

Software Engineer (2009–2017)

Search Consultant · Lucidworks India | Senior Software Engineer · Target.com India | Sr. Java Programmer · Happiest Minds | Search Consultant · Sony India | Software Engineer · PointCross Life Sciences | Search Intern · Glassdoor USA